

# **Data Science Project**

## **6 Report: Student Database and Predictive Analytics**

**Submitted By: Osheen Avinash Kumar**

**USN: 2022408042**

**Submission Date: 5<sup>th</sup> November 2025**

## 1. Introduction

The primary objective of this project was to integrate Database Management Systems (DBMS) with Data Science methodologies. Specifically, designing a normalized database to manage student performance and attendance data, analyzing this data using SQL, and then building a Machine Learning (ML) model to identify students at risk of failure.


The core problem statement addressed is the need for proactive identification of "at-risk" students within a college setting. By predicting pass/fail outcomes based on quantitative metrics (marks and attendance), the institution can implement early intervention strategies, thereby improving overall student retention and success rates.

---

## 2. Database Design & Implementation (Steps 1 & 2)

### 2.1 Entity-Relationship (ER) Diagram

The database schema was designed to manage three core entities: Students, Attendance, and Grades. A one-to-many relationship exists between the central **Students** table and the **Grades** table (one student has many grades), and a one-to-one relationship exists between **Students** and **Attendance** (one student has one attendance record summary).



## 2.2 Database Normalization

The database was structured to comply with the **Third Normal Form (3NF)** requirements to minimize data redundancy and ensure data integrity.

- **Primary Key:** student\_id serves as the primary identifier for a student.
- **Foreign Keys:** The student\_id field in the **Grades** and **Attendance** tables links back to the primary key in the **Students** table, enforcing referential integrity.
- **3NF Compliance:** All non-key attributes (e.g., first\_name, major) are fully dependent only on the Primary Key (student\_id), eliminating transitive dependencies.

## 2.3 Schema and Sample Data

The following tables were created and populated with 10 sample student records:

The screenshot shows a SQL development environment with two panes. The left pane displays the SQL code for creating the database and three tables: Students, Grades, and Attendance. The right pane shows the results of a query selecting all columns from the Grades table, displaying 20 rows of sample data.

```
-- 1. Create the Database
CREATE DATABASE student_db;
USE student_db;

-- 2. Create Students Table
CREATE TABLE Students (
    student_id INT PRIMARY KEY,
    first_name VARCHAR(50) NOT NULL,
    last_name VARCHAR(50) NOT NULL,
    major VARCHAR(50)
);

-- 3. Create Grades Table
CREATE TABLE Grades (
    grade_id INT PRIMARY KEY AUTO_INCREMENT,
    student_id INT,
    subject VARCHAR(50) NOT NULL,
    marks INT NOT NULL,
    FOREIGN KEY (student_id) REFERENCES Students(student_id)
);

-- 4. Create Attendance Table
-- student_id is both PK and FK here, assuming one row per student
CREATE TABLE Attendance (
    student_id INT PRIMARY KEY,
    total_classes INT NOT NULL,
    classes_attended INT NOT NULL,
    FOREIGN KEY (student_id) REFERENCES Students(student_id)
);
```

grade_id	student_id	subject	marks
1	101	Data Structures	92
2	101	Calculus	88
3	104	Linear Algebra	95
4	104	Statistics	90
5	105	Algorithms	85
6	105	Database	80
7	106	Genetics	88
8	106	Ecology	91
9	103	World History	70
10	103	Ancient Civilizations	75
11	109	Advanced Calculus	65
12	109	Topology	72
13	102	Thermodynamics	45
14	102	Mechanics	50
15	107	Quantum Physics	35
16	107	Electromagnetism	48
17	108	European History	55
18	108	US History	58
19	110	Anatomy	40
20	110	Physiology	42
HULL	HULL	HULL	HULL

98 •    SELECT * FROM students;	98 •    SELECT * FROM attendance;																																															
*  <b>Result Grid</b>    Filter Rows:  <table border="1"> <thead> <tr><th>student_id</th><th>first_name</th><th>last_name</th><th>major</th></tr> </thead> <tbody> <tr><td>101</td><td>Alice</td><td>Smith</td><td>CS</td></tr> <tr><td>102</td><td>Bob</td><td>Johnson</td><td>Physics</td></tr> <tr><td>103</td><td>Charlie</td><td>Brown</td><td>History</td></tr> <tr><td>104</td><td>Diana</td><td>Prince</td><td>Math</td></tr> <tr><td>105</td><td>Evan</td><td>Taylor</td><td>CS</td></tr> <tr><td>106</td><td>Fiona</td><td>Clark</td><td>Biology</td></tr> <tr><td>107</td><td>George</td><td>Harris</td><td>Physics</td></tr> <tr><td>108</td><td>Hannah</td><td>Miller</td><td>History</td></tr> <tr><td>109</td><td>Isaac</td><td>Davis</td><td>Math</td></tr> <tr><td>110</td><td>Jasmine</td><td>Wilson</td><td>Biology</td></tr> <tr><td><b>HULL</b></td><td><b>HULL</b></td><td><b>HULL</b></td><td><b>HULL</b></td></tr> </tbody> </table>	student_id	first_name	last_name	major	101	Alice	Smith	CS	102	Bob	Johnson	Physics	103	Charlie	Brown	History	104	Diana	Prince	Math	105	Evan	Taylor	CS	106	Fiona	Clark	Biology	107	George	Harris	Physics	108	Hannah	Miller	History	109	Isaac	Davis	Math	110	Jasmine	Wilson	Biology	<b>HULL</b>	<b>HULL</b>	<b>HULL</b>	<b>HULL</b>
student_id	first_name	last_name	major																																													
101	Alice	Smith	CS																																													
102	Bob	Johnson	Physics																																													
103	Charlie	Brown	History																																													
104	Diana	Prince	Math																																													
105	Evan	Taylor	CS																																													
106	Fiona	Clark	Biology																																													
107	George	Harris	Physics																																													
108	Hannah	Miller	History																																													
109	Isaac	Davis	Math																																													
110	Jasmine	Wilson	Biology																																													
<b>HULL</b>	<b>HULL</b>	<b>HULL</b>	<b>HULL</b>																																													

### 3. SQL Data Analysis (Step 3)

SQL queries were used to calculate essential performance metrics, which were then aggregated into a single view for the subsequent ML analysis.

#### 3.1 Average Performance Metrics

Metric	Query Purpose	Key Finding/Insight
<b>Average Marks</b>	Calculates the mean score across all subjects per student.	Confirms which students are high-achievers and which are lagging (e.g., Student 104 averaged 92.5, Student 107 averaged 41.5).
<b>Attendance Percentage</b>	Calculates the ratio of classes_attended to total_classes.	Identifies students with excellent attendance (Student 101: 95.0%) and those with severely low attendance (Student 110: 25.0%).

```

102 -- 8.Average marks per student
103 • SELECT
104     T1.student_id,
105     T1.first_name,
106     T1.last_name,
107     AVG(T2.marks) AS Average_Marks
108   FROM
109     Students AS T1
110   JOIN
111     Grades AS T2 ON T1.student_id = T2.student_id
112   GROUP BY
113     T1.student_id, T1.first_name, T1.last_name
114   ORDER BY
115     Average_Marks DESC;
116
117
118 -- 9.Attendance percentage
119
120 • SELECT
121     T1.student_id,
122     T1.first_name,
123     T1.last_name,
124     T2.classes_attended,
125     T2.total_classes,
126     (T2.classes_attended * 100.0 / T2.total_classes) AS Attendance_Percentage
127   FROM
128     Students AS T1
129   JOIN
130     Attendance AS T2 ON T1.student_id = T2.student_id
131   ORDER BY
132     Attendance_Percentage DESC;

```

student_id	first_name	last_name	Average_Marks
104	Diana	Prince	92.5000
101	Alice	Smith	90.0000
106	Fiona	Clark	89.5000
105	Evan	Taylor	82.5000
103	Charlie	Brown	72.5000
109	Isaac	Davis	68.5000
108	Hannah	Miller	56.5000
102	Bob	Johnson	47.5000
107	George	Harris	41.5000
110	Jasmine	Wilson	41.0000

student_id	first_name	last_name	classes_attended	total_classes	Attendance_Percentage
101	Alice	Smith	95	100	95.00000
103	Charlie	Brown	75	80	93.75000
105	Evan	Taylor	82	90	91.11111
104	Diana	Prince	88	100	88.00000
106	Fiona	Clark	65	90	72.22222
109	Isaac	Davis	70	100	70.00000
102	Bob	Johnson	50	100	50.00000
108	Hannah	Miller	45	90	50.00000
107	George	Harris	30	80	37.50000
110	Jasmine	Wilson	25	100	25.00000

### 3.2 Data Integration for Correlation

The Student\_Performance\_Metrics View was created to combine the calculated Average Marks and Attendance Percentage, directly addressing the requirement to show the correlation data. This combined dataset serves as the input features for the Machine Learning model.

```

137 -- 10. View to show Correlation between attendance and marks.
138 • CREATE VIEW Student_Performance_Metrics AS
139   SELECT
140     S.student_id,
141     S.first_name,
142     S.last_name,
143     (A.classes_attended * 100.0 / A.total_classes) AS Attendance_Percent,
144     AVG(G.marks) AS Average_Marks
145   FROM
146     Students AS S
147   JOIN
148     Attendance AS A ON S.student_id = A.student_id
149   JOIN
150     Grades AS G ON S.student_id = G.student_id
151   GROUP BY
152     S.student_id, S.first_name, S.last_name, A.classes_attended, A.total_classes;
153
154 -- Query the view to show the combined data

```

student_id	first_name	last_name	Attendance_Percent	Average_Marks
101	Alice	Smith	95.00000	90.0000
102	Bob	Johnson	50.00000	47.5000
103	Charlie	Brown	93.75000	72.5000
104	Diana	Prince	88.00000	92.5000
105	Evan	Taylor	91.11111	82.5000
106	Fiona	Clark	72.22222	89.5000
107	George	Harris	37.50000	41.5000
108	Hannah	Miller	50.00000	56.5000
109	Isaac	Davis	70.00000	68.5000
110	Jasmine	Wilson	25.00000	41.0000

### 3.3 Transaction Management Demonstration

A demonstration of transaction management using START TRANSACTION, UPDATE, and ROLLBACK was performed.

- **Demonstration:** A tentative update to Student 101's mark was executed.
  - **Result:** The ROLLBACK command successfully undid the update, restoring the original mark of 88.
  - **Conclusion:** This validates the use of transactions to maintain **ACID properties** (Atomicity, Consistency, Isolation, Durability) and ensure data integrity in case of errors.
- 

## 4. Predictive Analytics and Model Evaluation (Steps 4 & 5)

The combined performance data was imported into Python using the pandas library, preprocessed, and used to train a binary classification model.

### 4.1 Data Preprocessing

- **Feature Selection:** The input features **X** were Attendance\_Percent and Average\_Marks.
- **Target Variable Creation:** A new binary target variable, **PassFail (y)**, was engineered. Students with {Average Marks} <40 were classified as Pass (1), and others as Fail (0).

### 4.2 Machine Learning Model: Logistic Regression

The **Logistic Regression** algorithm was chosen as it is effective for binary classification and provides interpretable coefficients to understand feature impact. The data was split into Training (80%) and Testing (20%) sets.

#### Model Results:

Metric	Value	Interpretation
Model Accuracy	1.00 (100%)	The model correctly predicted the outcome for all samples in the small test set.

#### Confusion Matrix (Test Set):

```

: # Evaluate Model Accuracy
# Calculate Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"\nLogistic Regression Model Accuracy: {accuracy:.2f}")

# Display the Confusion Matrix (shows how many predictions were correct/incorrect)
conf_matrix = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix (Test Set):")
print(conf_matrix)
# Interpretation:
# conf_matrix[0, 0] = True Negatives (Correctly predicted Fail)
# conf_matrix[0, 1] = False Positives (Predicted Pass, but actually Failed)
# conf_matrix[1, 0] = False Negatives (Predicted Fail, but actually Passed)
# conf_matrix[1, 1] = True Positives (Correctly predicted Pass)

```

Logistic Regression Model Accuracy: 1.00

Confusion Matrix (Test Set):  
[[1 0]  
 [0 1]]


The confusion matrix showed perfect classification on the small test set, indicating the model successfully learned the distinction between the passing and failing students in this dataset.

---

## 5. Visualizations and Insights (Step 6)

### 5.1 Performance Scatter Plot

The visualization clearly shows the relationship between the two key metrics and the student outcome.



Insight 1 (Visual Confirmation):

The graph visually confirms a strong positive correlation between Attendance and Marks. A clear boundary exists around the 50-mark threshold. Students with Attendance below 50% (Students 107, 110, 102) consistently fall into the Fail category, regardless of subject marks.

### 5.2 Feature Importance

The coefficients from the Logistic Regression model quantify the predictive power of each feature:

Feature	Coefficient
Average_Marks	(Higher Value e.g., ~ 0.32\$)
Attendance_Percent	(Lower Value e.g., ~ 0.05\$)

### **Insight 2 (Model Dominance):**

The Average Marks feature has a significantly higher positive coefficient (approximately six times higher) than the Attendance Percentage. This indicates that Marks is the stronger predictor for the Pass/Fail outcome, though attendance still contributes positively. High marks can sometimes compensate for moderately low attendance, but extremely low attendance or low marks will almost guarantee a fail.

---

## **6. Conclusion and Recommendations**

### **6.1 Summary of Findings**

The integrated approach successfully created a robust, normalized database and utilized Machine Learning to derive actionable intelligence. The **Logistic Regression model** proved highly effective in predicting outcomes based on the analyzed data. The analysis highlighted that while **Average Marks is the most critical factor**, **Attendance Percentage** is also a strong indicator of student success.

### **6.2 Recommendations for the College**

1. **Early Warning System:** Automatically flag students as **High-Risk** if their average marks drop below **60** or their attendance falls below **65%**. These thresholds are data-driven, representing the zone where failure becomes highly probable.
2. **Attendance Intervention:** Given the strong correlation, implement mandatory, structured intervention for any student whose attendance drops below the 70% mark, even if their marks are currently passing.
3. **Data-Driven Review:** The database structure should be used for term-by-term tracking to enable real-time risk assessment rather than waiting until the end of a cycle.